



Welcome to the **Co**Grammar Data Cleaning

The session will start shortly...

Questions? Drop them in the chat. We'll have dedicated moderators answering questions.



Data Science Session Housekeeping

- The use of disrespectful language is prohibited in the questions, this is a supportive, learning environment for all - please engage accordingly.
(Fundamental British Values: Mutual Respect and Tolerance)
- No question is daft or silly - **ask them!**
- There are **Q&A sessions** midway and at the end of the session, should you wish to ask any follow-up questions. Moderators are going to be answering questions as the session progresses as well.
- If you have any questions outside of this lecture, or that are not answered during this lecture, please do submit these for upcoming Academic Sessions. You can submit these questions here: [Questions](#)

Data Science Session Housekeeping cont.

- For all **non-academic questions**, please submit a query:
www.hyperiondev.com/support
- Report a **safeguarding** incident:
www.hyperiondev.com/safeguardreporting
- We would love your **feedback** on lectures: [Feedback on Lectures](#)

CoGrammar

Data Cleaning

April 2024

Data Cleaning

- ❖ Data cleaning is a crucial step in the data science pipeline
- ❖ Ensures **data quality and reliability** for analysis and modeling
- ❖ Common data quality issues include **missing data, duplicates, inconsistent formatting, and outliers**

Learning objectives

- ❖ Describe techniques for **handling missing data** and when each is appropriate to use.
- ❖ Demonstrate how to identify and remove duplicate records in a dataset using Pandas.
- ❖ Explain the importance of **consistent data formatting and apply methods to standardise data.**
- ❖ Define **outliers** and discuss strategies for detecting and handling them appropriately based on the data context.

Missing Data



Handling Missing Data

- ❖ Missing data refers to the **absence of values in one or more variables in a dataset.**
- ❖ Identifying missing values:
 - Look for **null, NaN, or empty cells in the dataset.**
 - Use functions like **isnull()** or **isna()** in Pandas

```
# Identify missing values
```

```
df_missing.isnull().sum()
```

```
✓ 0.0s
```

```
total_bill    10
```

```
tip           0
```

```
sex          10
```

```
smoker       0
```

```
day          0
```

```
time        0
```

```
size         0
```

```
dtype: int64
```




Understand Missing Data Mechanisms



- ❖ These mechanisms are more important if you do research in the field, so we're going to glaze over it
- ❖ It helps us understand what techniques to use, but **we could intuit it in most cases**
- ❖ Here is the [original paper](#)

Understand Missing Data Mechanisms

- ❖ **MCAR:** Missing Completely at Random (missingness unrelated to any variables)
 - Smoking status is not recorded in a random sample of patients
- ❖ **MAR:** Missing at Random (missingness depends on observed variables)
 - Smoking status is not documented in female patients because the doctor was too shy to ask 😊
- ❖ **MNAR:** Missing Not at Random (missingness depends on missing values themselves)
 - Smoking status is not recorded in patients admitted as an emergency, who are also more likely to have worse outcomes from surgery

Techniques for Handling Missing Data

- ❖ **Deletion:** Remove records with missing values (only suitable if missing data is minimal and random).
 - Suitable for random missingness
 - Not the first resort, dropping data means losing some important context or skewing the dataset in some cases

```
df.shape
```

```
✓ 0.0s
```

```
(244, 7)
```

```
# Deletion: Remove records with missing values
```

```
df_deleted = df_missing.dropna()
```

```
df_deleted.shape
```

```
✓ 0.0s
```

```
(225, 7)
```

Techniques for Handling Missing Data

❖ **Imputation:** Fill in missing values with estimated or calculated values.

➤ Simple imputation: Mean, median, or mode imputation

```
# Simple Imputation: Fill missing values with mean for numeric columns and mode for categorical  
# columns  
df_imputed = df_missing.copy()  
df_imputed['total_bill'] = df_imputed['total_bill'].fillna(df_imputed['total_bill'].mean())  
df_imputed['sex'] = df_imputed['sex'].fillna(df_imputed['sex'].mode()[0])
```

➤ Advanced imputation: K-Nearest Neighbors (KNN), Multiple Imputation by Chained Equations (MICE)

■ We'll get to KNN in another lecture 😊

Techniques for Handling Missing Data

```
# Advanced Imputation: KNN Imputation
from sklearn.impute import KNNImputer

# Create a copy of the dataset for KNN imputation
df_imputed_knn = df_missing.copy()

# Initialize and fit the KNN imputer
imputer = KNNImputer(n_neighbors=5)
df_imputed_knn[['total_bill', 'tip', 'size']] = imputer.fit_transform(
    df_imputed_knn[['total_bill', 'tip', 'size']]
)
```

Duplicates



Dealing with Duplicates

- ❖ Identify duplicates using functions like `duplicated()` in Pandas

```
# Show all duplicated rows  
df_duplicates[df_duplicates.duplicated(keep=False)]
```

`keep = False` just marks all duplicates.

	total_bill	tip	sex	smoker	day	time	size
46	22.23	5.00	Male	No	Sun	Dinner	2
92	5.75	1.00	Female	Yes	Fri	Dinner	2
123	15.95	2.00	Male	No	Thur	Lunch	2
158	13.39	2.61	Female	No	Sun	Dinner	2
198	13.00	2.00	Female	Yes	Thur	Lunch	2
202	13.00	2.00	Female	Yes	Thur	Lunch	2
234	15.53	3.00	Male	Yes	Sat	Dinner	2
244	22.23	5.00	Male	No	Sun	Dinner	2
245	15.53	3.00	Male	Yes	Sat	Dinner	2
246	13.39	2.61	Female	No	Sun	Dinner	2
247	5.75	1.00	Female	Yes	Fri	Dinner	2
248	15.95	2.00	Male	No	Thur	Lunch	2

Dealing with Duplicates

- ❖ Dropping duplicates is fine and encouraged, it does not cause the data to lost necessary context

```
# Remove duplicate records  
df_deduplicated = df_duplicates.drop_duplicates()
```

Data Formatting and Standardization

- ❖ Consistent data formatting is essential for accurate analysis and compatibility
- ❖ Common formatting issues:
 - **Date and time formats:** Ensure consistent representation (e.g., YYYY-MM-DD, HH:MM:SS)
 - **Text case inconsistencies:** Convert text to a consistent case (lowercase or uppercase)
 - **Inconsistent value representations:** Standardize values (e.g., "Yes"/"No" vs. "Y"/"N")

Data Formatting and Standardization

- ❖ Techniques for standardizing data:
 - Convert date/time columns using `to_datetime()`
 - Convert text case using `str.lower()` or `str.upper()`
 - Map inconsistent values to standardized representations

```
# Standardize text case for 'sex' and 'smoker' columns
df['sex'] = df['sex'].str.upper()
df['smoker'] = df['smoker'].str.title()
```

```
df.head()
```

✓ 0.0s

	total_bill	tip	sex	smoker	day	time	size
0	16.99	1.01	FEMALE	No	Sun	Dinner	2
1	10.34	1.66	MALE	No	Sun	Dinner	3
2	21.01	3.50	MALE	No	Sun	Dinner	3
3	23.68	3.31	MALE	No	Sun	Dinner	2
4	24.59	3.61	FEMALE	No	Sun	Dinner	4

Outliers





Outliers

- ❖ Outliers are data points that significantly deviate from the rest of the data distribution



Identifying Outliers

- ❖ Visual inspection using plots like box plots, scatter plots, or histograms



Identifying Outliers

- ❖ Statistical methods like z-score or interquartile range (IQR)
 - Much less common given how good box plot already show outliers

```
# Identify outliers using z-score
from scipy import stats

z_scores = np.abs(stats.zscore(df['total_bill']))
threshold = 2.5
outliers_zscore = np.where(z_scores > threshold)

outliers_zscore
```

✓ 0.0s

```
(array([ 59, 102, 156, 170, 182, 197, 212]),)
```

```
# Identify outliers using Interquartile Range (IQR)
Q1 = df['total_bill'].quantile(0.25)
Q3 = df['total_bill'].quantile(0.75)
IQR = Q3 - Q1

outliers_iqr = df[(df['total_bill'] < (Q1 - 1.5 * IQR)) | (df['total_bill'] > (Q3 + 1.5 * IQR))]
len(outliers_iqr)
```

✓ 0.0s

MagicPython

9

Handling Outliers

- ❖ **Removal:** Remove outliers if they are erroneous or irrelevant to the analysis
 - Use when outliers are clearly erroneous or irrelevant to the analysis
 - Be cautious, as removing outliers may result in loss of information

```
# Removal: Remove outliers
```

```
df_removed = df[~((df['total_bill'] < (Q1 - 1.5 * IQR)) | (df['total_bill'] > (Q3 + 1.5 * IQR)))]
```

```
df_removed.shape
```

```
✓ 0.0s
```


```
MagicPython
```

```
(235, 9)
```

Handling Outliers

- ❖ **Transformation:** Apply mathematical transformations (e.g., logarithmic, square root) to reduce the impact of outliers
 - Use when outliers are valid but have a skewed distribution
 - Helps to reduce the impact of outliers while preserving the data


Handling Outliers



Handling Outliers

- ❖ **Winsorization:** Replace extreme values with the nearest non-outlier values
 - Use when outliers are valid but need to be treated to reduce their influence
 - Maintains the overall distribution shape while limiting the extreme values

Handling Outliers



Iterative Data Cleaning



Iterating

- ❖ Data cleaning is an iterative process that may require multiple rounds
- ❖ Continuously assess and refine the cleaned data based on analysis results and feedback
- ❖ Integrate data cleaning with data analysis and modeling for optimal results

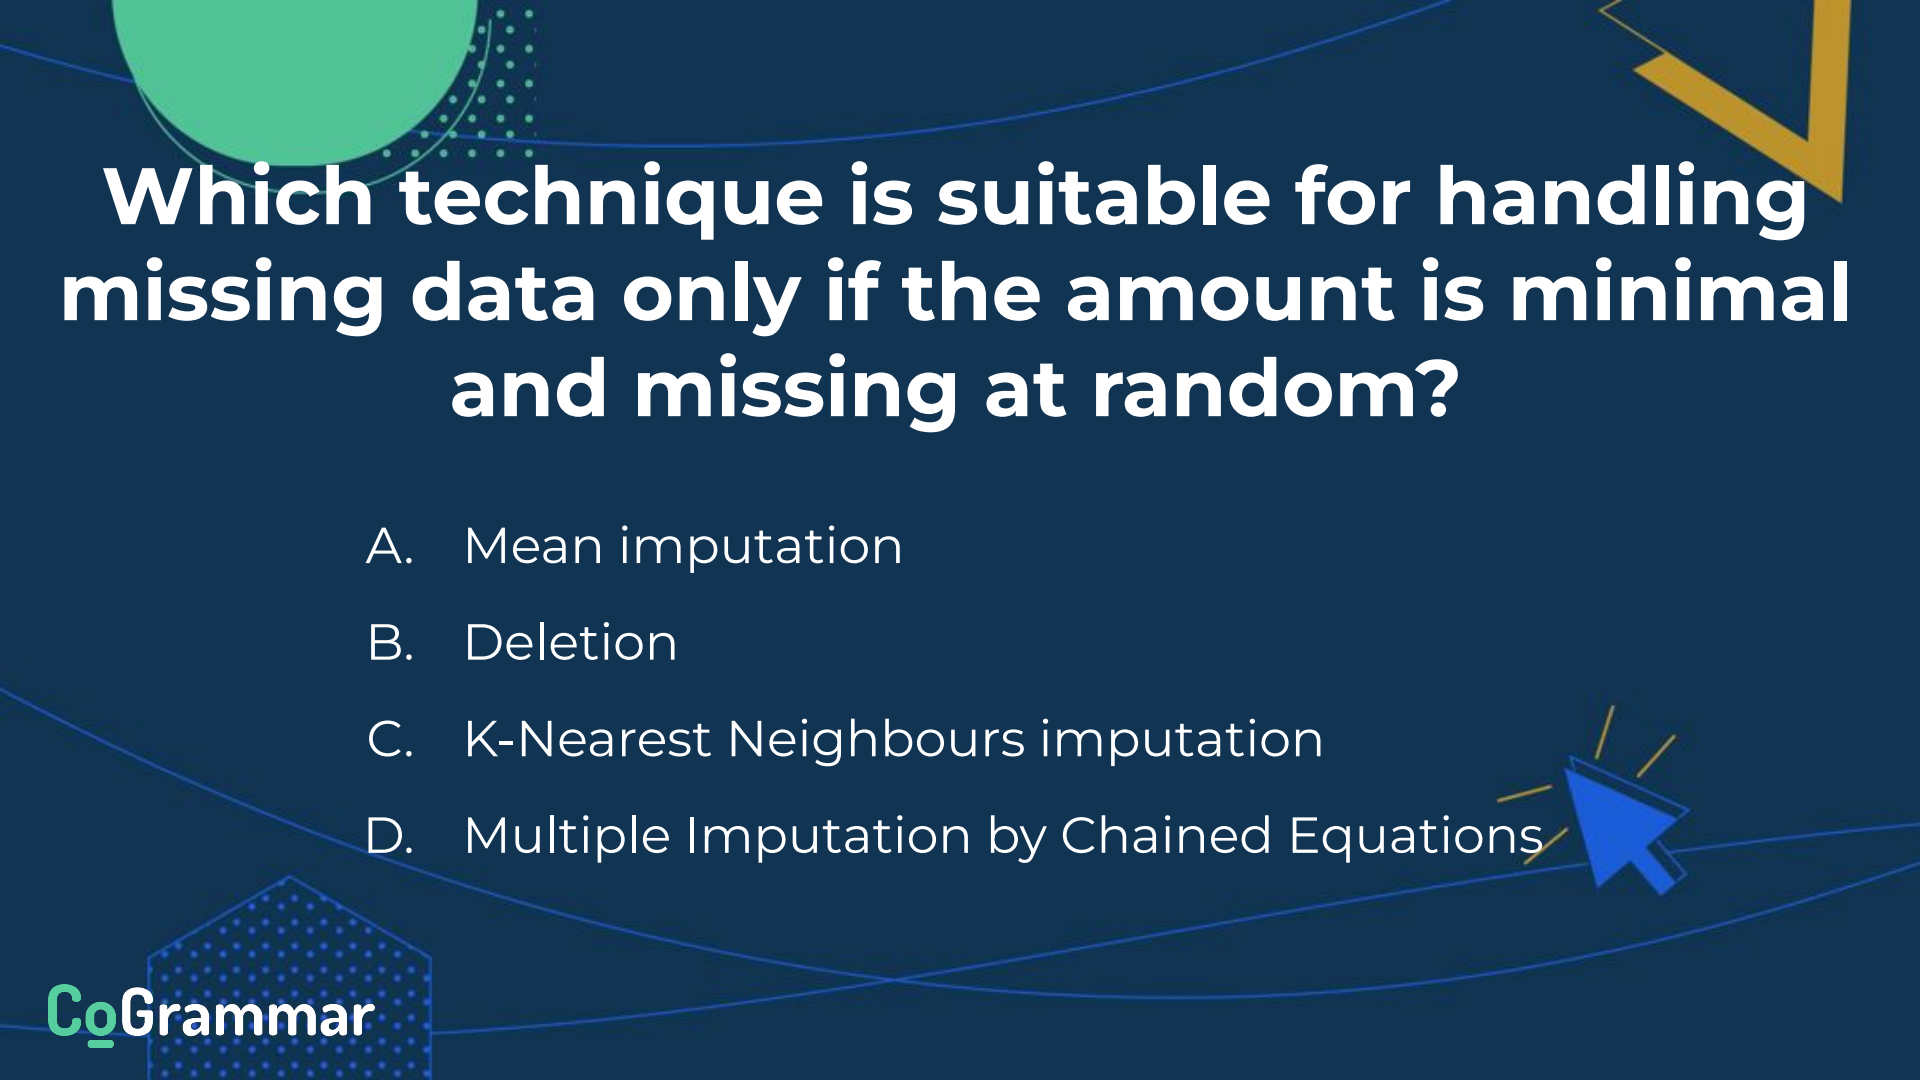


Which of the following is NOT a common data quality issue?

- A. Missing values
- B. Duplicates
- C. Inconsistent formatting
- D. Small sample size

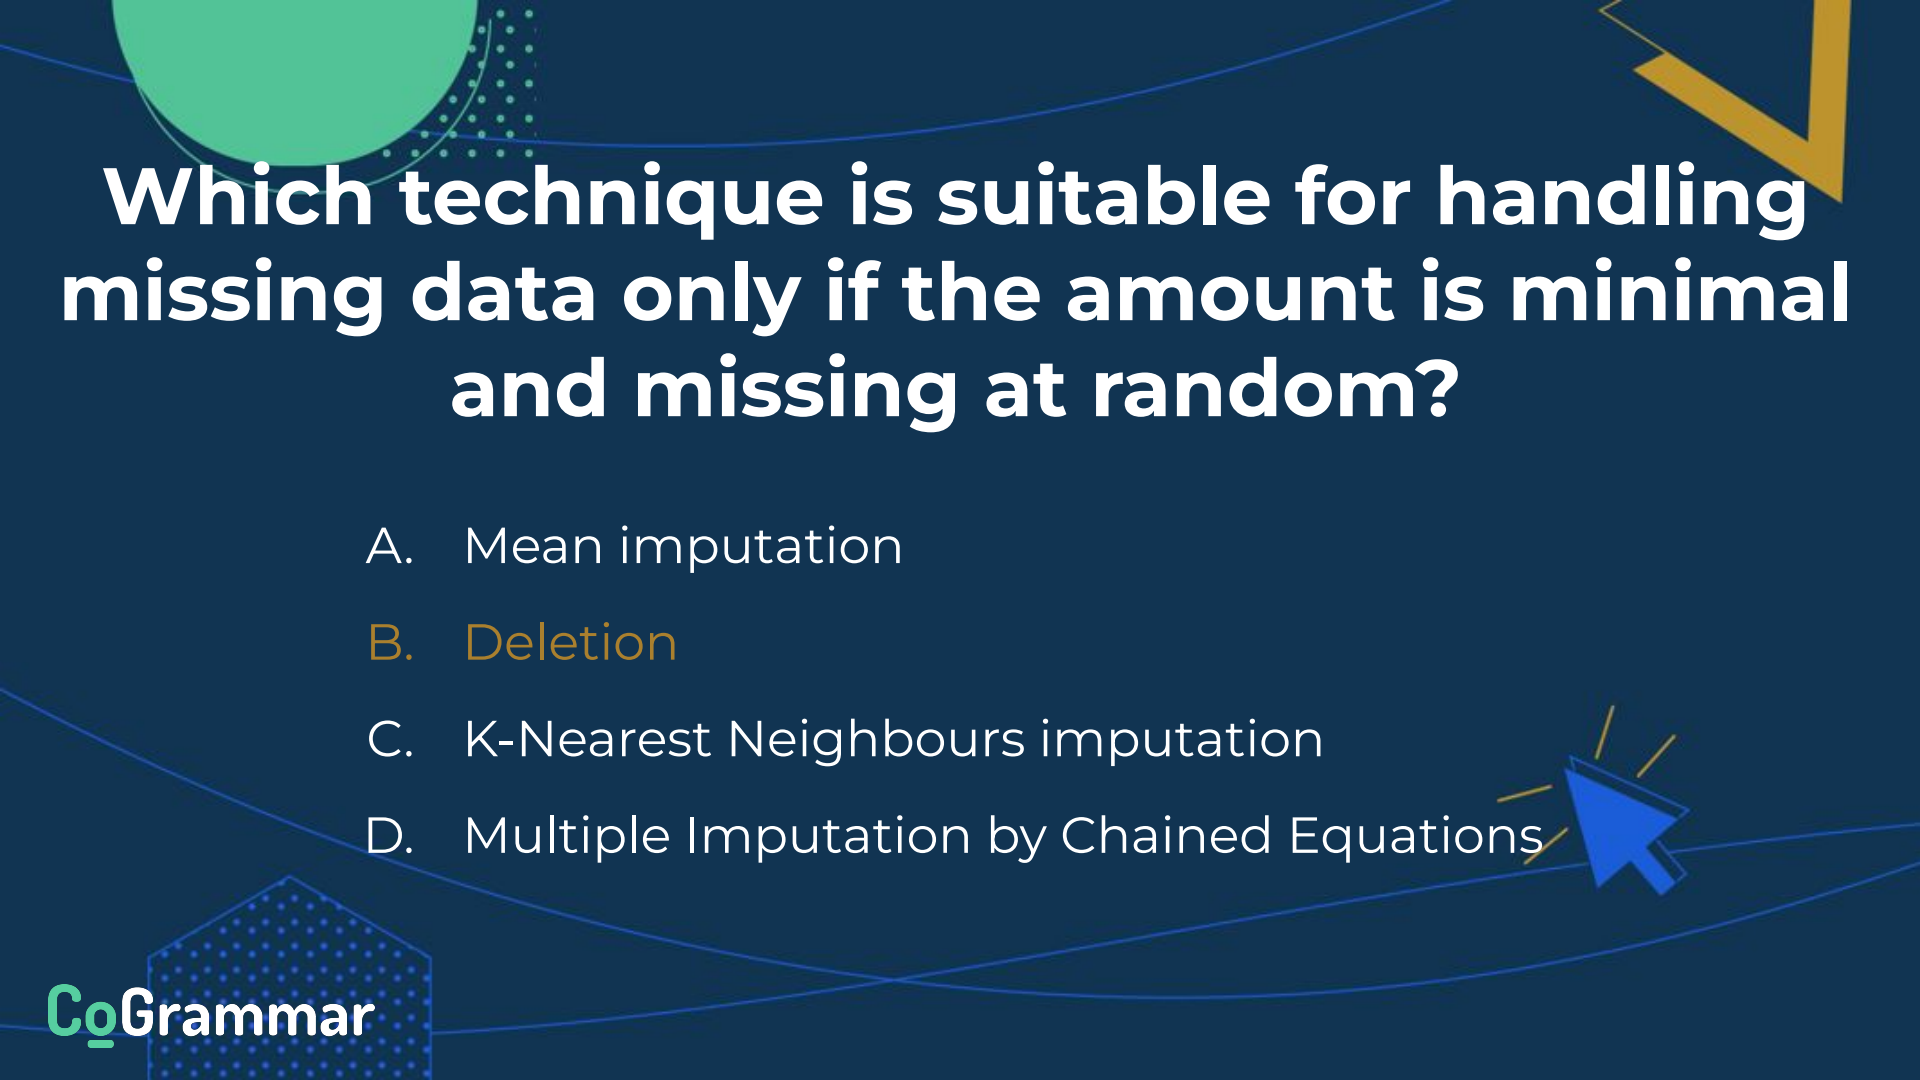
Which of the following is NOT a common data quality issue?

- A. Missing values
- B. Duplicates
- C. Inconsistent formatting
- D. Small sample size



Which technique is suitable for handling missing data only if the amount is minimal and missing at random?

- A. Mean imputation
- B. Deletion
- C. K-Nearest Neighbours imputation
- D. Multiple Imputation by Chained Equations



Which technique is suitable for handling missing data only if the amount is minimal and missing at random?

- A. Mean imputation
- B. Deletion
- C. K-Nearest Neighbours imputation
- D. Multiple Imputation by Chained Equations



In Pandas, which function can be used to identify duplicate records in a dataset?

- A. `find_duplicates()`
- B. `uplicated()`
- C. `is_duplicate()`
- D. `has_duplicates()`



In Pandas, which function can be used to identify duplicate records in a dataset?

- A. `find_duplicates()`
- B. `uplicated()`
- C. `is_duplicate()`
- D. `has_duplicates()`



Which of the following is a technique for standardising inconsistent text case?

- A. `astype()`
- B. `to_datetime()`
- C. `upper()` or `lower()`
- D. `strip()`

Which of the following is a technique for standardising inconsistent text case?

- A. `astype()`
- B. `to_datetime()`
- C. `upper()` or `lower()`
- D. `strip()`

Which strategy replaces outlier values with the nearest non-outlier values?

- A. Removal
- B. Transformation
- C. Winsorisation
- D. Standardisation

Which strategy replaces outlier values with the nearest non-outlier values?

- A. Removal
- B. Transformation
- C. Winsorisation
- D. Standardisation



Further Learning

KDNuggets - [Learn Data Cleaning and Preprocessing for Data Science with This Free eBook](#)

Kaggle - [Short Data Cleaning Course](#)

Questions and Answers



Thank you for attending



Department
for Education

CoGrammar

